

# TRPO

Trust Region Policy Optimization

# Advantage

- Advantage (A):  $A(s, a) = Q(s, a) - V(s)$
- Intuitive: How good an action is compared to the average action for a specific state.

- TRPO is an on-policy algorithm.
- TRPO updates policies by taking the largest step possible to improve performance
- It does this while satisfying a special constraint on how close the new and old policies are allowed to be.
- The constraint is expressed in terms of **KL-Divergence**, a measure of (something like, but not exactly) distance between probability distributions.

# Trust Region

Well known method in optimization

You have a function to optimize and you have a local approximation of this function - which works in a boundary called the trust region

But it gets really inaccurate if you go away from your starting point

Trust Region - Space where we trust our local approximation

Making sure that any action that is performed is well within the region that we Trust to be a good action.

- This is different from normal policy gradient, which keeps new and old policies close in parameter space.
- But even seemingly small differences in parameter space can have very large differences in performance—so a single bad step can collapse the policy performance.
- This makes it dangerous to use large step sizes with vanilla policy gradients, thus hurting its sample efficiency.
- TRPO nicely avoids this kind of collapse, and tends to improve performance.

```

) import pandas as pd
import yfinance as yf

class YahooDownloader:
    """Provides methods for retrieving daily stock data from
    Yahoo Finance API
    Attributes
    -----
        start_date : str
            start date of the data (modified from config.py)
        end_date : str
            end date of the data (modified from config.py)
        ticker_list : list
            a list of stock tickers (modified from config.py)
    Methods
    -----
        fetch_data()
            Fetches data from yahoo API
    """

```

## Compared to A2C

- We can use a modified version of the loss function given below:

$$L_{\theta_{old}} = \mathbb{E}_t \left[ \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)} A_t \right]$$

- Instead of the loss function in A2C:

$$L(\theta) = \mathbb{E}_t \left[ \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A_t \right]$$

Let  $\pi_\theta$  denote a policy with parameters  $\theta$ . The theoretical TRPO update is:

$$\begin{aligned}\theta_{k+1} &= \arg \max_{\theta} \mathcal{L}(\theta_k, \theta) \\ \text{s.t. } \bar{D}_{KL}(\theta || \theta_k) &\leq \delta\end{aligned}$$

where  $\mathcal{L}(\theta_k, \theta)$  is the *surrogate advantage*, a measure of how policy  $\pi_\theta$  performs relative to the old policy  $\pi_{\theta_k}$  using data from the old policy:

$$\mathcal{L}(\theta_k, \theta) = \mathbb{E}_{s, a \sim \pi_{\theta_k}} \left[ \frac{\pi_\theta(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right],$$

and  $\bar{D}_{KL}(\theta || \theta_k)$  is an average KL-divergence between policies across states visited by the old policy:

$$\bar{D}_{KL}(\theta || \theta_k) = \mathbb{E}_{s \sim \pi_{\theta_k}} [D_{KL}(\pi_\theta(\cdot|s) || \pi_{\theta_k}(\cdot|s))].$$